

CSL 301 - Operating System

XV6-public Lab: TLB and Page Fault Measurement

Objective

Understand TLB behavior, page faults, and lazy allocation in `xv6-public` by modifying kernel and user code. You will build incrementally with guided steps.

Part 1: Add a Page Fault Counter in `proc` and Create Syscall `getpagefaults()`

Step 1: Add a counter field `page_faults` to struct `proc`. File to modify: `proc.h`

```
// Add an integer field in struct proc named page_faults.  
// #WRITE YOUR CODE HERE
```

Step 2: Initialize it in `allocproc()` (`proc.c`). Write the line that sets page faults counter initially to zero.

```
// #WRITE YOUR CODE HERE
```

Step 3: Implement the syscall to fetch page faults count.

- Assign a syscall number for `getpagefaults` in `syscall.h`.
- Declare the syscall handler in `sysproc.c`. You may use the following hint to implement it:

```
int sys_getpagefaults(void) {  
    // Hint: return the page_faults field of current process.  
    // #WRITE YOUR CODE HERE  
}
```

- Register your syscall in `syscall.c` and add user-space prototypes in `user.h`.
- Add the syscall stub in `usys.S`.

Part 2: Lazy Page Allocation

Step 1: Modify the `vmfault` function to allocate pages lazily. The full code is provided below; study it carefully and understand its working:

```
int vmfault(pde_t *pgdir, uint va, int write) {
    struct proc *p = myproc();
    char *mem;
    if (va >= p->sz)
        return -1;
    va = PGROUNDDOWN(va);
    if (walkpgdir(pgdir, (void *)va, 0))
        return 0;
    mem = kalloc();
    if (mem == 0)
        return -1;
    memset(mem, 0, PGSIZE);
    if (mappages(pgdir, (void *)va, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0)
    {
        kfree(mem);
        return -1;
    }
    return 0;
}
```

Step 2: Modify the trap handler to increase page faults count and handle faults by calling `vmfault()`. Replace the relevant code section with:

```
// Inside trap() for page faults
p->page_faults++;
if (vmfault(p->pgdir, rcr2(), tf->err & 2) < 0)
    p->killed = 1;
```

Part 3: User Programs to Measure Page Faults

Task: Complete the user program `tlbrun.c` to: - allocate increasing number of pages - access pages many times to generate faults - print page counts, trials, ticks, and faults

Fill the missing code marked below:

```

#include "types.h"
#include "user.h"

#define PAGESIZE 4096
#define MAXPAGES 1024

int main() {
    int jump = PAGESIZE/sizeof(int);
    printf(1, 'PageCount\tTrials\tTicks\tPageFaults\n');
    for (int numpages = 1; numpages <= MAXPAGES; numpages *= 2) {
        int trials = 5000000;
        int faults_before = getpagefaults();
        int start = uptime();
        int *arr = (int*) sbrk(numpages * PAGESIZE);
        if (arr == (void*) -1)
            exit();
        for (int t = 0; t < trials; t++) {
            for (int i = 0; i < (numpages/2)*jump; i += jump) {
                // #WRITE YOUR CODE HERE: Access the page
                to trigger faults
            }
        }
        int end = uptime();
        int faults_after = getpagefaults();
        printf(1, '%d\t%d\t%d\t%d\n', numpages, trials, end-start
, faults_after - faults_before);
    }
    exit();
}

```

Similarly, complete `tlbtest.c` to accept page count and trials from command line and print results. Use the above as guide.

Part 4: Integration and Testing

Add the programs `_tlbrun` and `_tlbtest` to the `UPROGS` in `Makefile`. Build and run tests.

Submission Checklist

- Kernel source changes with comments.
- User programs (`tlbrun.c`, `tlbtest.c`) with explanations of your approach.
- Sample output demonstrating page fault counts and timing.